

CS 111 Midterm exam

CHU; JONATHAN WEI-HONG

TOTAL POINTS

63 / 100

QUESTION 1

1 Dirty bits 3 / 10

- 0 pts Correct

- 10 pts No answer

- 5 pts Incorrect in explaining why dirty bit improves performance

✓ - 5 pts Incorrect in explaining how dirty bit improves performance

✓ - 2 pts Did not link performance increase to less disk I/O

QUESTION 2

2 ABI and system call interface 7 / 10

- 0 pts The subset relationship is clearly stated in the answer.

- 10 pts No answer

✓ - 3 pts Wrote down some sentences related to question, but didn't clearly mention system call interface is a subset of ABI.

- 1 pts More close to the answer but still missing clearly mentioning the subset relationship.

QUESTION 3

3 Shared libraries 6 / 10

- 0 pts Correct

- 10 pts No answer

✓ - 4 pts Missing details: multiple processes accessing the shared global data at the same time would be a problem.

QUESTION 4

4 System calls and trap instructions 8 / 10

- 0 pts Correct

- 10 pts No answer

- 2 pts Did not mention transition of processor from

unprivileged mode to privileged mode

- 2 pts Did not mention OS runs appropriate code for the system call

- 2 pts Did not explain usage of trap handler

- 2 pts Did not mention OS will determine what trap was caused by trap instruction

✓ - 2 pts Did not mention associated parameters preset by user application are saved

QUESTION 5

5 Working sets and page stealing 10 / 10

✓ - 0 pts Correct

- 10 pts No answer

- 5 pts Incorrect description of working set.

- 5 pts Incorrect description of page stealing algorithms.

- 3 pts insufficient description of page stealing algorithms.

- 3 pts Insufficient description of working sets.

- 2 pts Important element is that page stealing takes pages from processes whose working sets are too large and gives them to processes whose working sets are too small.

- 3 pts Goal of a working set is not to maximize the number of pages in memory, but to figure out the right number to have there.

- 2 pts Working set has nothing to do with TLB.

- 1 pts Just because a working set is large doesn't mean it isn't using its pages.

- 2 pts working set is not really about preventing thrashing, since that can occur even with properly implemented working sets.

- 3 pts Important to note that working sets are associated with processes and are controlled by their behavior.

- 3 pts Page stealing is used to build proper

working sets, not vice versa.

- 2 pts Processes don't voluntarily release page frames. That's why it's called stealing.

- 2 pts Page table usually bigger than the working set.

- 0 pts Click here to replace this description.

QUESTION 6

6 Scheduling algorithm metrics 7 / 10

- 0 pts Correct

- 10 pts No answer

- 2 pts Maximizing jobs completed does not necessarily translate to maximum throughput, by most definitions.

- 8 pts Fairness not guaranteed by non-preemptive scheduling. Starvation not necessarily avoided, either.

- 1 pts SJF has nothing to do with number of pages.

- 5 pts Insufficient explanation of why metric is maximized.

✓ - 3 pts Not for all non-preemptive algorithms.

Turnaround time won't be optimized for non-preemptive FIFO, for example. Question asked about non-preemptive algorithms in general, not just one example of such an algorithm.

- 5 pts Turnaround time is not necessarily optimized. It's time of job arrival till time of job completion. With non-preemptive scheduling, one long-running job can kill the turnaround time of many other jobs.

- 0 pts As stated in the test instructions, nothing on the back of the page is graded.

- 5 pts Won't necessarily optimize time to completion. A long running job will not be interrupted, causing other short jobs to incur long times to completion. If you interrupted the long job for the short ones, average time to completion would improve.

- 10 pts Did not specify a metric.

- 4 pts "Minimizing context switches" isn't a performance metric, though doing so is likely to improve some metrics.

- 3 pts Insufficient explanation of why metric is maximized.

- 10 pts Response time may not be optimized with non-preemptive scheduling, since one long-running job can kill the response time of many other jobs.

- 2 pts Won't also optimize average time to completion, since long-running job can kill time to completion of many other jobs.

- 4 pts Throughput is typically defined as the amount of work produced by a system, not the number of jobs it completes. By the latter definition, non-preemptive scheduling doesn't optimize the metric, since you could finish many short jobs in the time it takes to finish one long job.

- 5 pts Not clear exactly what you mean by "process speed".

- 4 pts That's not the definition of mean response time. It's the average time to get some response from the system, not time to completion.

- 2 pts Not very clear description of chosen metric.

- 4 pts Non-preemptive scheduling is not likely to optimize number of deadlines met, since newly arrived jobs with short deadlines can't preempt a running job with a long deadline.

- 0 pts Not the usual definition of time to completion, but correct as described.

- 10 pts Round robin is not a non-preemptive algorithm

- 10 pts "operations/second->output" makes no sense. Output is not a metric.

- 6 pts Your assumptions are rarely true, and if not true, average time to completion is not optimized, by most definitions of that metric.

- 8 pts Incorrect description of throughput.

Throughput is not the same as turnaround time, and turnaround time is not necessarily optimized by non-preemptive scheduling.

- 3 pts Metric you're looking for is throughput, not "turnout" or turnaround time.

- 5 pts Fairness not guaranteed (by any definition) for all types of non-preemptive scheduling, such as non-preemptive shortest job first.

- 2 pts You're thinking of throughput, not total execution time.

- 10 pts Round robin is not a metric, it's a scheduling algorithm, and not even a preemptive one.

- 0 pts Click here to replace this description.

QUESTION 7

7 Worst fit and fragmentation 10 / 10

✓ - 0 pts Correct

- 3 pts Worst fit algorithms fit allocation requests into the largest free chunk of memory available, assuming no perfect fit is available. The remainder of that chunk will be as large as possible, meaning it will be well suited to match later requests.

- 3 pts A best fit algorithm will choose the free chunk closest and larger in size to the requested allocation, which implies that the leftover free memory returned to the free list is likely to be a small chunk, poorly suited to matching future requests.

- 4 pts The definition of external fragmentation is scattering small, useless chunks of free memory throughout the free list, so best fit is more likely to cause external fragmentation than worst fit.

- 10 pts wrong answer

- 10 pts No answer

QUESTION 8

8 Page tables for fork vs. shared memory

IPC 5 / 10

- 0 pts Correct

- 10 pts No answer

✓ - 5 pts Major difference is fork results in copy-on-write, while shared memory IPC doesn't.

- 1 pts new process has its own page table, but its contents are the same.

- 3 pts No discussion of fork page table issues.

- 1 pts Stack isn't shared in shared memory IPC.

- 2 pts Fork need not be followed by exec, leading to COW issues.

- 1 pts IPC shared memory almost always read/write, at least by one of the processes.

- 2 pts Data segment also likely to change after fork.

- 0 pts Not well worded, but I think you have the right idea.

- 10 pts So what's the difference?

- 10 pts What is the difference in their page table behavior?

- 3 pts Difference won't be in the TLB.

- 2 pts Copy-on-write issue.

- 5 pts Not a thread issue. Copy on write is the main relevant mechanism.

- 4 pts Shared memory doesn't share page tables. Just entries in different page tables point to the same page frame.

- 3 pts IPC is not about libraries, it's about data.

QUESTION 9

9 Condition variables 3 / 10

- 0 pts Correct

- 10 pts No answer

✓ - 4 pts A condition variable is used to determine if some specific pre-defined condition has or has not occurred.

- 3 pts If the condition does occur, one or more of the blocked processes will be unblocked and permitted to run.

✓ - 3 pts The condition variable allows a process to wait for a specific condition without requiring the process to use a busy loop to check for the condition's occurrence.

- 10 pts wrong answer

QUESTION 10

10 TLB misses 4 / 10

- 0 pts Correct

- 10 pts No answer

✓ - 3 pts Missing case of invalid entry.

- 4 pts Missing case of valid entry on disk.

- 3 pts Missing case of valid entry in RAM.

- 1 pts Page fault is on non-present, not invalid.

- 3 pts Case with page on disk is present bit not set. Invalid bit is different.

- 4 pts Different cases for valid page on disk and

valid page in RAM.

- 2 pts What happens for an invalid entry?

✓ - 1 pts **TLB is a cache of page table entries, not pages.**

- 1 pts Per test instructions, text on the back of the page is not graded.

- 2 pts First step is to consult in-RAM page table.

- 1 pts Disk isn't searched, since page table contains disk location of non-present pages.

- 2 pts More details on not present case.

- 3 pts Spatial locality does not play into TLB miss handling.

- 3 pts Memory won't be searched. Either the page is present, not present, or not valid. Present pages have their PTE loaded, not present pages are fetched from disk, invalid pages cause an exception.

- 2 pts Page table entry itself will indicate if page is on disk. No need to invoke clock algorithm.

- 1 pts Dirty bit doesn't indicate whether a page is in memory or not. Present bit does. Dirty bit indicates if an in-memory page has been written.

- 1 pts Invalid case is not that the page cannot be found, but that its PTE is marked invalid.

- 1 pts How is it determined if a segmentation fault should occur?

✓ - 2 pts **Not present pages are in the page table. They're just marked as "not present."**

Midterm Exam
CS 111, Principles of Operating Systems
Winter 2018

Name: Jonathan Chu

Student ID Number: 009852220

This is a closed book, closed note test. Answer all questions.

Each question should be answered in 2-5 sentences. DO NOT simply write everything you remember about the topic of the question. Answer the question that was asked. Extraneous information not related to the answer to the question will not improve your grade and may make it difficult to determine if the pertinent part of your answer is correct. Confine your answers to the space directly below each question. Only text in this space will be graded. No question requires a longer answer than the space provided.

1. In a virtual memory system, why is it beneficial to have a dirty bit associated with a page?

The dirty bit indicates whether a page has been modified. This is useful in VM systems because some systems will prefer to evict "clean", or unmodified, pages as opposed to "dirty" pages when running page replacement algorithms because modified pages might contain information important to the user.

2. What is the relationship between a system's Application Binary Interface and its system call interface?

A system's application binary interface sits "below" the system call interface. System calls allow user-mode processes to invoke the kernel, where the system's ABI handles low-level transitions of compiled binary programs and instructions.

3. Why can't shared libraries include global data?

Shared libraries are loaded upon linkage, so there is no way for these libraries to know what a program's global data is ahead of time, let alone use it in execution. Programs that use shared libraries must know ahead of time what functions and parameters to call from those libraries. In other words, shared libraries and the programs that use them are separate entities until they are linked, so global data cannot exist between them beforehand.

4. Describe how a trap instruction is used to implement a system call in a typical operating system.

A user-mode process may execute a system call to invoke the kernel and run privileged instructions. A trap instruction is the way the operating system knows to run kernel code. The trap table is consulted and tells the OS where in memory to jump to in order to execute the OS code associated with a particular type of trap which carries out the system call's intent.

5. What is the relationship between the concept of working sets and page stealing algorithms?

A process' ^{desired} working set can be thought of as its ideal size of allocated memory (the # of pages given to it) such that it may run efficiently with few page faults yet not use up too much space.

Page stealing algorithms are run upon page faults, when a process can't find a particular page, or in other words, they are run when a process' current working set is not quite satisfying its needs and page faults are occurring. The page stealing algorithm decides which process to take pages from in order to balance working set sizes across processes.

6. Name a performance metric that is likely to be maximizable using non-preemptive scheduling. Why is this form of scheduling useful to maximize this metric?

The average time to completion may be maximizable using non-preemptive scheduling assuming the scheduler knows to schedule shorter jobs first and assuming that there are multiple jobs waiting to be scheduled at a given time. This form of scheduling can maximize the average time to completion metric because in order to ensure a process completes as quickly as possible, it should run to completion without preemption.

We may also note that fairness and average time to first run are certainly not optimized by a non-preemptive scheduler because processes can hog the CPU.
and will

7. Why does a worst fit algorithm for managing a free list handle external fragmentation better than a best fit algorithm?

A worst fit algorithm handles external fragmentation better because it tends to create small unusable fragments more slowly than best fit by using the largest free segment each time. Best fit will instead use the smallest possible segment, leaving a greater likelihood of that segment being reduced to a small, unusable segment. The idea behind this is that using a larger free segment will leave a larger free segment behind once needed space is allocated.

8. Both shared memory IPC and the processes' data areas after a Linux fork() operation would require the page tables of two processes to point to the same physical page frames. What would be different about the two cases (other than being caused by IPC vs. forking)?

Shared memory IPC simply requires processes to have access to the same portion of physical memory. Everything else about the processes' states (address space, registers, PC) could be different. After a fork(), however, the original process' state is copied such that the two processes may continue execution of the same code with the same address space and registers.

9. What is the purpose of a condition variable?

A condition variable is used to halt a thread's execution until some condition is satisfied and its execution may continue. This is useful for threads running instructions that are dependent on the state or completion of say another thread. For example, a parent thread may want to wait for its child to finish execution before exiting. This would be implemented in C using `pthread_join()`.

10. In a system using demand paging, what operations are required when a TLB miss occurs? What are the possible outcomes of those operations?

In a system using demand paging, a TLB miss might often occur because the system does not load all necessary pages when it begins execution of a process but rather loads pages as they are needed. When a TLB miss occurs, the system consults the process' page table to look for the specified page. If the page is found, it is added to the TLB (and another page is removed, if the TLB was full) so that it may be loaded faster next time. If the page is not found in the page table, we have a page fault, and the page we're looking for is thus found and added to the page table, with the help of a page replacement algorithm if necessary.