

CS 111 Midterm

Jack Zhang

TOTAL POINTS

82 / 100

QUESTION 1

11 8 / 8

- ✓ - **0 pts** Correct
- **8 pts** No answer
- **7 pts** Wrong answer
- **4 pts** Answer on right track but not correct
- **3 pts** Answer needs more detail

QUESTION 2

2 2 7 / 10

- **0 pts** Correct
- **10 pts** No answer
- **9 pts** Wrong answer
- **3 pts** Incorrect answers for RR
- **3 pts** Incorrect answers for FCFS
- **3 pts** Incorrect answers for SJF
- ✓ - **3 pts** Answer of which has the largest overhead is incorrect or not present

☞ Round Robin is likely to have the largest overhead due to expensive context switching / being preemptive.

QUESTION 3

3 3 10 / 10

- ✓ - **0 pts** Correct
- **10 pts** No answer
- **9 pts** Wrong answer
- **5 pts** Answer on the right track but not correct OR missing part
- **3 pts** Answer needs a little more detail OR is slightly off

QUESTION 4

4 4 8 / 8

- ✓ - **0 pts** Correct

- **2 pts** Miss some details or some sentences are not accurate/correct enough.

- **5 pts** Wrote down something, but far from correct/enough.

- **7 pts** Wrong answer.

- **7 pts** Cannot fully understand/recognize your answer. Please type down your answer using regrading request. Thanks.

- **8 pts** No answer.

QUESTION 5

5 5 8 / 8

- ✓ - **0 pts** Correct
- **3 pts** Didn't explain for shared memory IPC, different processes refer to the exact same page frames or need synchronization.
- **3 pts** Didn't explain the copy-on-write property for fork.
- **6 pts** Wrong answer or not what we want.
- **7 pts** Cannot fully understand/recognize your answer. Please type down your answer using regrading request. Thanks.
- **8 pts** No answer.
- **3 pts** Missing details.

QUESTION 6

6 6 10 / 10

- ✓ - **0 pts** Correct
- **3 pts** Didn't consider the case where the page is in RAM.
- **3 pts** Didn't consider the case where the page is not in RAM but in disk (page fault).
- **6 pts** Wrote down something that makes sense, but didn't cover the main points that we are looking for. For example, didn't answer what operations are required (page table lookup) and didn't cover all

outcomes.

- **9 pts** Cannot fully understand/recognize your answer. Please type down your answer using regrading request. Thanks.
- **10 pts** No answer.
- **3 pts** Missing details.

QUESTION 7

7 7 15 / 15

- ✓ - **0 pts** Correct
- **5 pts** The first 4 iterations are page faults
- **2 pts** Missing last page fault
- **15 pts** Incorrect
- **10 pts** All squares were not filled out
- **5 pts** Incorrect use of the algorithm

QUESTION 8

8 8 15 / 15

- ✓ - **0 pts** Correct
- **15 pts** Incorrect/ Not Done
- **5 pts** Used bit should be set on load
- **5 pts** Page fault on startup
- **5 pts** Incorrect use of the algorithm
- **2 pts** Missing page fault

QUESTION 9

9 16 pts

9.1 a 1 / 4

- ✓ - **3 pts** Problematic
 - **4 pts** Incorrect
 - **0 pts** Correct
 - **2 pts** Partially correct
- ☞ We need to support the windows load module and emulate system calls.

9.2 b 0 / 3

- ✓ - **3 pts** Incorrect
 - **2 pts** Problematic
 - **0 pts** Click here to replace this description.
 - **1 pts** Partially correct
- ☞ A new 2nd level trap handler would be written

to intercept the Windows system calls, and pass it on to an emulation layer, which would try to simulate the effects of each Window's system call, using Solaris mechanisms.

9.3 c 0 / 3

- **1 pts** Partially correct.
 - **2 pts** Problematic
 - ✓ - **3 pts** Incorrect
 - **0 pts** Correct
- ☞ Performance should be okay since user-level instructions don't need to be emulated. Only system calls do.

9.4 d 0 / 3

- **2 pts** Problematic
 - **0 pts** Correct
 - ✓ - **3 pts** Incorrect
 - **1 pts** Partially correct
- ☞ The architecture does not change!

9.5 e 0 / 3

- **0 pts** Correct
- ✓ - **3 pts** Incorrect
- **2 pts** Click here to replace this description.

Midterm Examination
CS 111, Spring 2019
5/1/2019, 4 - 5:50pm

Name: Jack Zhang Student ID: 004993345

This is a closed book, closed notes test. One single-sided cheat sheet is allowed.

1. What is the benefit of using the copy-on-right optimization when performing a fork in the Linux system?

Instead of loading every copying and loading everything at once into the child process, we perform those operations ~~when it is needed~~ on demand when it is needed. This way, there is not ~~much~~ any overhead when spawning a child process through fork. Also saves space, potentially.

and the child process can read from the parent's pages until it actually writes to it, at which point we make a copy.

2. Round Robin, First come First Serve, and Shortest Job First are three scheduling algorithms that can be used to schedule a CPU. What are their advantages and disadvantages? Which one is likely to have the largest overhead? Why?

Round Robin promotes fairness by minimizing response time, thereby boosting interactivity. However, it is terrible for turnaround time. First come first serve, or FIFO is simple to implement but is bad for turnaround time when a time-consuming process job arrives first. It is also not good for response time b/c it does not preempt long jobs for ones that arrive. Shortest job first is better in terms of turnaround time than FIFO, but does not preempt, so is worse than "shortest time to completion". It also isn't good for response time. FIFO is most likely to have the most overhead in both turnaround time and response time if long jobs get queued first.

3. In a virtual memory system, why is it beneficial to have a dirty bit associated with a page? What are the techniques we can use to reduce the I/O involved in evicting dirty pages?

If a page in memory has been modified, then when it is evicted to the hard disk by the page-replacement daemon, it must actually write to the disk instead of just simply adding to the swap space. This is far more costly, so we use a dirty bit in the PTE to solve the problem. Building on the LRU approximation clock algorithm, which at first evicted the first PTE w/ use bit = 0, now we give priority to use bit 0 AND dirty bit = 0 (not modified), THEN use bit = 0 and dirty bit = 1. This way, costlier write operations are held off until they must be performed. We can also cluster the dirty pages together, only evicting them and writing to the disk in groups, b/c big writes are faster than many small writes.

4. What is the relationship between the concept of working sets and page stealing algorithms?

Working sets work by finding the optimal number of pages to allocate to a process, and evicting it if there is too many and page stealing if there is not enough.

from another process.

For fork, when the child writes to a page, ~~the~~ the VPN of that page no longer points to the PFN that the parent also points to. Instead, it makes a copy, which is a new value at a different physical address and holds a different value because it was just written to. Shared memory IPC, the VPNs of the threads (in) are constantly mapped to the same PFN, so a change in one thread is reflected in the other.

5. Both shared memory IPC and the processes' data areas after a Linux fork operation would require the page tables of two processes to point to the same physical page frames. What would be different about the two cases (other than being caused by IPC vs. forking)?

~~For Shared memory IPC, when one thread/process write to a page, it is changed & the change reflects for all threads/processes with VPN mapping to that PFN. Thus, there is only one PTE. However, for forks, there is no memory sharing, as a forked process is its own, independent process. Thus, each forked process, along with the parent process, has own PTE with its VPN mapping to the same PFN.~~

3b. if page not in disk throw error!

6. In a system using demand paging, what operations are required when a TLB miss occurs? What are the possible outcomes of these operations?

Outcomes:

1. TLB fault triggers fault handler, permission level raised to Kernel
2. Searches for the PFN in cache; not found = page fault
3. Page fault handler reads in the page from hard disk at address specified where the PFN usually occupies in the PTE.
4. Page is added to cache
5. Returns to do entire memory reference again → TLB miss, add page to TLB from cache
6. 3rd memory reference = TLB hit.

7. **Optimal LRU.** Consider the reference string shown along the top of the following graphical structure. The system has four frames. Use the LRU algorithm to select pages for replacement. Place the page number in the proper frame. Mark when page faults occur in the bottom line of boxes. State how many page faults occur. The numbers across the top indicate the reference string.

| | | | | | | | | | | | | | | | | | |
|---------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | 0 | 1 | 3 | 6 | 2 | 4 | 5 | 2 | | 0 | 3 | 1 | 2 | 5 | 4 | 1 | 0 |
| Frame 0 | 0 | 0 | 0 | 0 | 2 | 2 | 2 | 2 | 6 | 4 | 0 | 5 | 0 | 3 | 1 | 2 | 5 |
| 1 | | | | | | | | | 4 | 2 | 5 | 0 | 3 | 1 | 2 | 5 | 4 |
| 2 | | | | | | | | | 2 | 5 | 0 | 3 | 1 | 2 | 5 | 4 | 1 |
| 3 | | | | | | | | | 5 | 0 | 3 | 1 | 2 | 5 | 4 | 1 | 0 |
| Fault ? | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | | ✓ |

0 0 0 0 1 3 6 6

1 1 3 3 6 2 4 4

3 3 6 2 4 5 2

6 2 4 5 2

↗

14 page

14 page faults

8. **Clock Algorithm.** The clock algorithm is an approximation of LRU based on using one use bit for each page. When a page is used its use bit is set to 1. We also use a pointer to the next victim, which is initialized to the first page/frame. When a page is loaded, it is

set to point to the next frame. The list of pages is considered as a circular queue. When a page is considered for replacement, the use bit for the next victim page is examined. If it is zero [that page is replaced] otherwise [the use bit is set to zero, the next victim pointer is advanced, and the process repeated until a page is found with a zero use bit].

Consider the reference string shown along the top of the following graphical structure. The system has four frames. Use the clock algorithm described in the previous paragraph. The narrow boxes to the right of the page number boxes can be used to keep up with use bits. Place the page number in the proper frame. Mark when page faults occur in the bottom line of boxes. State how many page faults occur.

12 page faults

| | | | | | | | | | | | | | | | | | |
|---------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | 0 | 1 | 3 | 6 | 2 | 4 | 5 | 2 | 5 | 0 | 3 | 1 | 2 | 5 | 4 | 1 | 0 |
| Frame 0 | 0 | 0 | 0 | 0 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 0 |
| 1 | | 1 | 1 | 1 | 1 | 4 | 4 | 4 | 4 | 4 | 3 | 3 | 3 | 3 | 4 | 4 | 4 |
| 2 | | | 3 | 3 | 3 | 3 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |
| 3 | | | | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 0 | 0 | 1 | 1 | 1 | 1 | 1 |
| Fault ? | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | | | ✓ | ✓ | ✓ | | | ✓ | | ✓ |

not used!
(blank)
filled a bit
accidentally

9. In the early 1990s, SUN Microsystems, the maker of the Solaris Operating System, wanted to move from the engineering desktop, where it was well established, to a broader market for personal productivity tools. The best personal productivity tools were all being written for Windows platforms, and SUN was on the wrong side of the applications/demand/volume cycle, which made getting those applications ported to Solaris a non-option.

One approach to their problem was to modify the version of Solaris that ran on x86 processors (the popular hardware platform for Windows) to be able to run Windows binaries without any alterations to those binaries. This would allow Sun to automatically offer all of the great applications that were available for Windows.

- (a) What would have to be done to permit Windows binaries to be loaded into memory and executed on a Solaris/x86 system?

machine-level and conventions
~~conform to Windows~~ ABI and ~~make a instructions~~ for calling
routines consistent w/ those of Windows, although
implementation may vary.

- (b) What would have to be done to correctly execute the system calls that the Windows programs requested?

Conform to ~~Windows~~ Windows API, modifying, for
example, the current Solaris OS implementam of read()
to take the same parameters, return the same values,
throw the same errors, etc. as ~~Windows~~ Window's read().

(c) How good might the performance of such a system be? Justify your answer.

It depends ~~on~~ on the implementation of the routines and subroutines themselves. If they are written to run more efficiently than Windows, then it would have good performance - it ~~doesn't~~ wouldn't matter at all that the programs were written for Windows.

(d) List another critical thing, besides supporting a new load module format and the basic system calls, that the system would have to be prepared to simulate? How might that be done?

The ~~API~~ API must be bound to the ISA.

Solaris follows its API and ABI

(e) Could a similar approach work on a Solaris/PowerPC or Solaris/SPARC system? Why or why not?

Yes, ~~the hardware~~ hardware platforms are similar.
if

Yes, as long as the ISA is written and ~~can~~ configured to conform to the Windows' ABI.